



# Docker Hadoop

# What is Docker?

- Lightweight VM (sort of correct)
- Good vehicle for delivering an application
- Runs on Linux, Mac and Windows
- Images versus containers

# Running Docker Hadoop

```
$ docker pull sequenceiq/hadoop-docker:2.3.0
```

```
/var/lib/docker  
~/Library/Containers/com.docker.docker  
~/docker/machine/machines/default
```

```
$ docker run -it sequenceiq/hadoop-docker:2.3.0 /etc/bootstrap.sh -bash
```

```
/  
Starting sshd: [ OK ]  
Starting namenodes on [localhost]  
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-root-namenode-b524750d0773.out  
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-root-datanode-b524750d0773.out  
Starting secondary namenodes [0.0.0.0]  
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-root-secondarynamenode-b524750d0773.out  
starting yarn daemons  
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn--resourcemanager-b524750d0773.out  
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-root-nodemanager-b524750d0773.out  
bash-4.1#
```

```
$ cd $HADOOP_PREFIX
$ export JAR=share/hadoop/mapreduce/hadoop-mapreduce-examples-2.3.0.jar
$ bin/hadoop jar $JAR wordcount input output
```

```
$ bin/hdfs dfs -ls
```

Found 2 items

drwxr-xr-x	- root supergroup	0	2014-10-23 08:14	input
drwxr-xr-x	- root supergroup	0	2016-09-13 01:29	output

```
$ bin/hdfs dfs -ls output
```

Found 2 items

-rw-r--r--	1 root supergroup	0	2016-09-13 01:29	output/_SUCCESS
-rw-r--r--	1 root supergroup	30038	2016-09-13 01:29	output/part-r-00000

```
$ bin/hdfs dfs -cat output/*
```

!= 3

"" 6

"\$HADOOP\_CLASSPATH" 1

# WordCount.java

```
package org.apache.hadoop.examples;

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;
```

# WordCount.java

```
public class WordCount {  
  
    public static class TokenizerMapper  
        extends Mapper<Object, Text, Text, IntWritable>{  
  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(Object key, Text value, Context context  
            ) throws IOException, InterruptedException {  
  
            StringTokenizer itr = new StringTokenizer(value.toString());  
            while (itr.hasMoreTokens()) {  
                word.set(itr.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
}
```

# WordCount.java

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {

    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

# WordCount.java

```
public static void main(String[] args) throws Exception {  
  
    Configuration conf = new Configuration();  
    String[] otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();  
    if (otherArgs.length != 2) {  
        System.err.println("Usage: wordcount <in> <out>");  
        System.exit(2);  
    }  
    Job job = new Job(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizerMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path(otherArgs[0]));  
    FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```



# Assignments

